

14_SVD

March 20, 2017

1 14 Linear Algebra: Singular Value Decomposition

One can always decompose a matrix A

$$A = U \operatorname{diag}(w_j) V^T \quad (1)$$

$$U^T U = U U^T = 1 \quad (2)$$

$$V^T V = V V^T = 1 \quad (3)$$

where the w_j are the *singular values*.

The inverse (if it exists) can be directly calculated from the SVD:

$$A^{-1} = V \operatorname{diag}(1/w_j) U^T$$

1.1 Solving ill-conditioned coupled linear equations

```
In [ ]: import numpy as np
```

1.1.1 Non-singular matrix

```
In [111]: A = np.array([
            [1, 2, 3],
            [3, 2, 1],
            [-1, -2, -6],
          ])
          b = np.array([0, 1, -1])
```

```
In [112]: np.linalg.solve(A, b)
```

```
Out[112]: array([ 0.83333333, -0.91666667,  0.33333333])
```

```
In [113]: U, w, VT = np.linalg.svd(A)
          print(w)
```

```
[ 7.74140616  2.96605874  0.52261473]
```

```
In [114]: U.dot(np.diag(w).dot(VT))
```

```

Out[114]: array([[ 1.,  2.,  3.],
                 [ 3.,  2.,  1.],
                 [-1., -2., -6.]])

In [115]: np.allclose(A, U.dot(np.diag(w).dot(VT)))

Out[115]: True

In [116]: inv_w = 1/w
           print(inv_w)

[ 0.1291755  0.33714774  1.91345545]

In [117]: A_inv = VT.T.dot(np.diag(inv_w)).dot(U.T)
           print(A_inv)

[[ -8.33333333e-01  5.00000000e-01 -3.33333333e-01]
 [  1.41666667e+00 -2.50000000e-01  6.66666667e-01]
 [ -3.33333333e-01 -1.11022302e-16 -3.33333333e-01]]

In [118]: np.allclose(A_inv, np.linalg.inv(A))

Out[118]: True

In [119]: x = A_inv.dot(b)
           print(x)
           np.allclose(x, np.linalg.solve(A, b))

[ 0.83333333 -0.91666667  0.33333333]

Out[119]: True

In [120]: A.dot(x)

Out[120]: array([-6.66133815e-16,  1.00000000e+00, -1.00000000e+00])

In [121]: np.allclose(A.dot(x), b)

Out[121]: True

```

1.1.2 Singular matrix

```

In [376]: C = np.array([
                [ 0.87119148,  0.9330127, -0.9330127],
                [ 1.1160254,  0.04736717, -0.04736717],
                [ 1.1160254,  0.04736717, -0.04736717],
            ])
           b1 = np.array([ 2.3674474, -0.24813392, -0.24813392])
           b2 = np.array([0, 1, 1])

```

```
In [310]: np.linalg.solve(C, b1)
```

```
-----  
  
LinAlgError                                Traceback (most recent call last)  
  
  <ipython-input-310-d689ff5cc60e> in <module>()  
----> 1 np.linalg.solve(C, b)  
  
/opt/local/Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/s  
382     signature = 'DD->D' if isComplexType(t) else 'dd->d'  
383     extobj = get_linalg_error_extobj(_raise_linalgerror_singular)  
--> 384     r = gufunc(a, b, signature=signature, extobj=extobj)  
385  
386     return wrap(r.astype(result_t, copy=False))  
  
/opt/local/Library/Frameworks/Python.framework/Versions/3.5/lib/python3.5/s  
88  
89 def _raise_linalgerror_singular(err, flag):  
--> 90     raise LinAlgError("Singular matrix")  
91  
92 def _raise_linalgerror_nonposdef(err, flag):
```

```
LinAlgError: Singular matrix
```

```
In [377]: U, w, VT = np.linalg.svd(C)  
          print(w)
```

```
[ 1.99999999e+00  1.00000000e+00  2.46519033e-32]
```

```
In [378]: U.dot(np.diag(w).dot(VT))
```

```
Out[378]: array([[ 0.87119148,  0.9330127 , -0.9330127 ],  
                 [ 1.1160254 ,  0.04736717, -0.04736717],  
                 [ 1.1160254 ,  0.04736717, -0.04736717]])
```

```
In [379]: np.allclose(C, U.dot(np.diag(w).dot(VT)))
```

```
Out[379]: True
```

```
In [380]: singular_values = np.abs(w) < 1e-12  
          print(singular_values)
```

```
[False False  True]
```

```
In [381]: inv_w = 1/w
          inv_w[singular_values] = 0
          print(inv_w)
```

```
[ 0.5  1.   0. ]
```

```
In [382]: C_inv = VT.T.dot(np.diag(inv_w)).dot(U.T)
          print(C_inv)
```

```
[[-0.04736717  0.46650635  0.46650635]
 [ 0.5580127  -0.21779787 -0.21779787]
 [-0.5580127   0.21779787  0.21779787]]
```

```
In [383]: x1 = C_inv.dot(b1)
          print(x1)
```

```
[-0.34365138  1.4291518  -1.4291518 ]
```

```
In [384]: C.dot(x1)
```

```
Out[384]: array([ 2.3674474 , -0.24813392, -0.24813392])
```

```
In [385]: C.dot(x1) - b1
```

```
Out[385]: array([-4.44089210e-16,  4.99600361e-16,  4.99600361e-16])
```

- The columns $U_{:,i}$ of U (i.e. $U.T[i]$ or $U[:, i]$) corresponding to non-zero w_i , i.e. $\{i : w_i \neq 0\}$, form the basis for the *range* of the matrix A .
- The columns $V_{:,i}$ of V (i.e. $V.T[i]$ or $V[:, i]$) corresponding to zero w_i , i.e. $\{i : w_i = 0\}$, form the basis for the *null space* of the matrix A .

Note that $x1$ can be written as a linear combination of $U.T[0]$ and $U.T[1]$:

```
In [408]: x1
```

```
Out[408]: array([-0.34365138,  1.4291518 , -1.4291518 ])
```

```
In [402]: U.T
```

```
Out[402]: array([[ -7.07106782e-01,  -4.99999999e-01,  -4.99999999e-01],
 [  7.07106780e-01,  -5.00000001e-01,  -5.00000001e-01],
 [ -2.47010760e-16,  -7.07106781e-01,   7.07106781e-01]])
```

```
In [410]: VT
```

```
Out[410]: array([[ -0.8660254 , -0.35355339,  0.35355339],
                [ -0.5       ,  0.61237244, -0.61237244],
                [ -0.        , -0.70710678, -0.70710678]])
```

```
In [411]: U.T[0].dot(x1), U.T[1].dot(x1)
```

```
Out[411]: (0.24299822382783731, -0.24299822305983199)
```

```
In [412]: VT[2].dot(x1)
```

```
Out[412]: 2.2204460492503131e-16
```

```
In [413]: U.T[0].dot(x1) * U.T[0] + U.T[1].dot(x1) * U.T[1] + 2 * VT[2]
```

```
Out[413]: array([-0.34365138, -1.41421356, -1.41421356])
```

The solution vector x_2 is in the null space:

```
In [349]: x2 = C_inv.dot(b2)
          print(x2)
          print(C.dot(x2))
          print(C.dot(x2) - b2)
```

```
[ 0.9330127  -0.43559574  0.43559574]
[ -3.33066907e-16  1.00000000e+00  1.00000000e+00]
[ -3.33066907e-16 -3.33066907e-16 -3.33066907e-16]
```

```
In [352]: C.dot(10*x2)
```

```
Out[352]: array([ -3.55271368e-15,  1.00000000e+01,  1.00000000e+01])
```

```
In [350]: C.dot(VT[2])
```

```
Out[350]: array([ 0.00000000e+00, -6.93889390e-18, -6.93889390e-18])
```

```
In [351]: VT[2]
```

```
Out[351]: array([-0.        , -0.70710678, -0.70710678])
```

```
In [138]: null_basis = VT[singular_values]
```

```
In [140]: C.dot(null_basis.T)
```

```
Out[140]: array([[ 0.00000000e+00],
                [ 4.44089210e-16],
                [ 0.00000000e+00]])
```

1.2 SVD for fewer equations than unknowns

M equations for N unknowns with $M < N$:

- no unique solutions (underdetermined)
- $N - M$ dimensional family of solutions
- SVD: at least $N - M$ zero or negligible w_j : columns of V corresponding to singular w_j span the solution space when added to a particular solution.

1.3 SVD for more equations than unknowns

M equations for N unknowns with $M > N$:

- no exact solutions in general (overdetermined)
- but: SVD can provide best solution in the least-square sense

$$\mathbf{x} = V \operatorname{diag}(1/w_j) U^T \mathbf{b}$$

where

- \mathbf{x} is a N -dimensional vector of the unknowns,
- V is a $N \times M$ matrix
- the w_j form a square $M \times M$ matrix,
- U is a $N \times M$ matrix (and U^T is a $M \times N$ matrix), and
- \mathbf{b} is the M -dimensional vector of the given values

It provides the \mathbf{x} that minimizes the residual

$$\mathbf{r} := |\mathbf{Ax} - \mathbf{b}|.$$

1.3.1 Linear least-squares fitting

This is the *linear least-squares fitting problem*: Given data points (x_i, y_i) , fit to a linear model $y(x)$, which can be any linear combination of functions of x .

For example:

$$y(x) = a_1 + a_2x + a_3x^2 + \dots + a_Mx^{M-1}$$

or in general

$$y(x) = \sum_{k=1}^M a_k X_k(x)$$

The goal is to determine the coefficients a_k .

Define the **merit function**

$$\chi^2 = \sum_{i=1}^N \left[\frac{y_i - \sum_{k=1}^M a_k X_k(x_i)}{\sigma_i} \right]^2$$

(sum of squared deviations, weighted with standard deviations σ_i on the y_i).

Best parameters a_k are the ones that *minimize* χ^2 .

Design matrix \mathbf{A} ($N \times M$, $N \geq M$), vector of measurements \mathbf{b} (N -dim) and parameter vector \mathbf{a} (M -dim):

$$A_{ij} = \frac{X_j(x_i)}{\sigma_i} \quad (4)$$

$$b_i = \frac{y_i}{\sigma_i} \quad (5)$$

$$\mathbf{a} = (a_1, a_2, \dots, a_M) \quad (6)$$

Minimum occurs when the derivative vanishes:

$$0 = \frac{\partial \chi^2}{\partial a_k} = \sum_{i=1}^N \sigma_i^{-2} \left[y_i - \sum_{k=1}^M a_k X_k(x_i) \right] X_k(x_i), \quad 1 \leq k \leq M$$

(M coupled equations)

$$\sum_{j=1}^M \alpha_{kj} a_j = \beta_k \quad (7)$$

$$\alpha \mathbf{a} = \beta \quad (8)$$

with the $M \times M$ matrix

$$\alpha_{kj} = \sum_{i=1}^N \frac{X_j(x_i) X_k(x_i)}{\sigma_i^2} \quad (9)$$

$$\alpha = \mathbf{A}^T \mathbf{A} \quad (10)$$

and the vector of length M

$$\beta_k = \sum_{i=1}^N \frac{y_i X_k(x_i)}{\sigma_i^2} \quad (11)$$

$$\beta = \mathbf{A}^T \mathbf{b} \quad (12)$$

The inverse of α is related to the uncertainties in the parameters:

$$\mathbf{C} := \alpha^{-1}$$

in particular

$$\sigma(a_i) = C_{ii}$$

(and the C_{ij} are the co-variances).

Solution of the linear least-squares fitting problem with SVD We need to solve the overdetermined system of M coupled equations

$$\sum_{j=1}^M \alpha_{kj} a_j = \beta_k \quad (13)$$

$$\alpha \mathbf{a} = \beta \quad (14)$$

SVD finds \mathbf{a} that minimizes

$$\chi^2 = |\mathbf{A}\mathbf{a} - \mathbf{b}|$$

The errors are

$$\sigma^2(a_j) = \sum_{i=1}^M \left(\frac{V_{ji}}{w_i} \right)^2$$

Example Synthetic data

$$y(x) = 3 \sin x - 2 \sin 3x + \sin 4x$$

with noise r added (uniform in range $-5 < r < 5$).

```
In [2]: import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
matplotlib.style.use('ggplot')


import numpy as np

In [3]: def signal(x, noise=0):
        r = np.random.uniform(-noise, noise, len(x))
        return 3*np.sin(x) - 2*np.sin(3*x) + np.sin(4*x) + r

In [6]: X = np.linspace(-10, 10, 500)
        Y = signal(X, noise=5)

In [7]: plt.plot(X, Y, 'r-', X, signal(X, noise=0), 'k--')

Out[7]: [<matplotlib.lines.Line2D at 0x10765da90>,
         <matplotlib.lines.Line2D at 0x10765dc50>]
```

```
In [8]: def fitfunc(x, a):
        return a[0]*np.cos(x) + a[1]*np.sin(x) + \
               a[2]*np.cos(2*x) + a[3]*np.sin(2*x) + \
               a[4]*np.cos(3*x) + a[5]*np.sin(3*x) + \
               a[6]*np.cos(4*x) + a[7]*np.sin(4*x)

        def basisfuncs(x):
            return np.array([np.cos(x), np.sin(x),
                             np.cos(2*x), np.sin(2*x),
                             np.cos(3*x), np.sin(3*x),
                             np.cos(4*x), np.sin(4*x)])

In [9]: M = 8
        sigma = 1.
        alpha = np.zeros((M, M))
        beta = np.zeros(M)
        for x in X:
            Xk = basisfuncs(x)
            for k in range(M):
                for j in range(M):
                    alpha[k, j] += Xk[k]*Xk[j]
        for x, y in zip(X, Y):
            beta += y * basisfuncs(x)/sigma

In [10]: U, w, VT = np.linalg.svd(alpha)
         V = VT.T
```

In this case, the singular values do not immediately show if any basis functions are superfluous (this would be the case for values close to 0).

```
In [11]: w
```

```
Out[11]: array([ 296.92809624,  282.94804954,  243.7895787 ,  235.7300808 ,
                235.15938555,  235.14838812,  235.14821093,  235.14821013])
```

... nevertheless, remember to routinely mask any singular values or close to singular values:

```
In [12]: w_inv = 1/w
         w_inv[np.abs(w) < 1e-12] = 0
         alpha_inv = V.dot(np.diag(w_inv)).dot(U.T)
```


Compare the fitted values to the original parameters $a_j = 0, +3, 0, 0, 0, -2, 0, +1$.

```
In [15]: a_values = alpha_inv.dot(beta)
         print(a_values)
```

```
[-0.01571487  3.15536126  0.14649695  0.27795436  0.04792114 -2.17431709
 -0.38617743  0.72017528]
```

```
In [14]: plt.plot(X, fitfunc(X, a_values), 'b-')
         plt.plot(X, signal(X, noise=0), 'k--')
```

```
Out[14]: [<matplotlib.lines.Line2D at 0x1075cbba8>]
```



```
In [ ]:
```